Machine Learning
3rd Lab Assignment
Instituto Superior Técnico

Neural Networks

This assignment aims at illustrating the applications of neural networks to image classification. We'll train a multilayer perceptron (MLP) and a convolutional neural network (CNN) for classification and compare the two approaches.

`Note` In this assignment, you'll often find between brackets, as in {`command`}, suggestions of Python commands that may be useful to perform the requested tasks. You should search Python documentation, when necessary, to obtain a description of how to use these commands.

1 Image Recognition

Our classification problem is an image recognition one, using supervised learning. Our goal is to classify images of clothing with 28x28 pixels each, into 10 categories. The following figure shows some of the images and their labels.
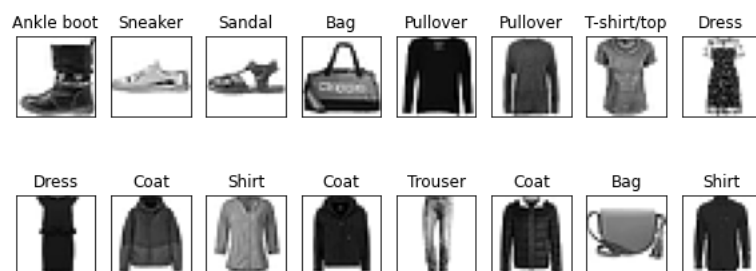

Figure 1: Example images.

1.1 Data

We will use the well known Fashion MNIST dataset which contains 70,000 grayscale images. The data set is already separated into 60000 training images and 10000 test images.
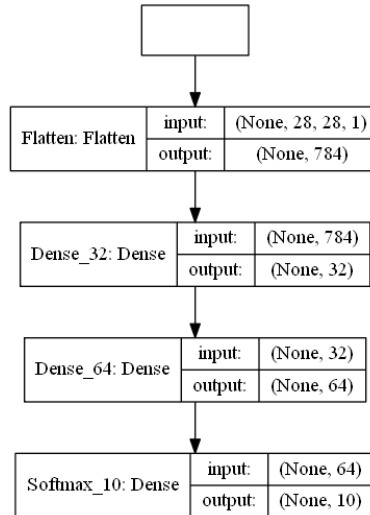
Figure 2: Graph plot of the MLP model.

The labels, which are between 0 and 9, correspond to the following categories: 'T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat','Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot'.

1. Load the data files and check the size of training and test data and corresponding labels.{`load_data from keras_datasets.fashion_mnist`}

2. Display some of the images in the train and test data. {`imshow, show from matplotlib`}

3. Divide data by 255 to get floating point values in the 0–1 range.

4. Convert train labels to one-hot encoding. In this representation, the label matrix will have 10 elements, with the component that corresponds to the pattern's class equal to 1, and all the other components equal to 0. {`to_categorical from keras`}

5. Split train data into two subsets, one for actual training and the other for validation. Use 20% for validation.{`train_test_split from sklearn`}

6. Add a 4^{th} dimension to your datasets so that the shape of patterns is (28,28,1) instead of (28,28){`expand_dims from numpy`}

1.2 MLP

Figure 2 shows a graph plot of the multi layer perceptron (MLP) you will create.

1. Create a sequential model and start by adding a flatten layer to convert the 2D images to 1D. Since this is the first layer in your model, you'll need to to specify the input

shape which, for 28x28 pixels with only one color channel (grayscale) is (28,28,1). {`Sequential, Add, Flatten from keras`}

2. Add two hidden layers to your MLP, the first with 32 and the second with 64 neurons. Use 'relu' as activation function for all neurons in the hidden layers.{`Add, Dense from keras`}

3. End the network with a softmax layer. This is a dense layer that will return an array of 10 probability scores, one per class, summing to 1. {`Dense from keras`}

4. Get the summary of your network to check it is correct.{`summary from keras`}

5. Create an Early stopping monitor that will stop training when the validation loss is not improving (use patience=10 and restore_best_weights=True). {`EarlyStopping from keras.callbacks`}

6. Fit the MLP to your training and validation data using 'categorical_crossentropy' as the loss function, a batch size of 200 and Adam as the optimizer (learning rate=0.001, clipnorm=1). Choose, as stopping criterion, the number of epochs reaching 200. Don't forget the Early Stopping callback.{`Compile, Fit from keras`}

7. Plot the evolution of the training loss and the validation loss. Note that the call to fit() returns a History object where metrics monitored during training are kept. {`Plot from matplotlib`}

8. To get an idea of how well the model generalizes to new, unseen data, evaluate performance (accuracy and confusion matrix) on the test data. {`predict from keras, accuracy_score, confusion_matrix from sklearn.metrics`}

9. Repeat the previous items without Early Stopping.

1.3 CNN

Figure 3 shows a graph plot of the Convolutional neural network you will create.

1. Create a Convolutional Neural Network (CNN) with two alternated Convolutional (with relu activation and 3x3 filters) and MaxPooling2D layers (2x2). Use 16 filters in the first conv layer and 16 in the second. Add a flatten layer and then a dense layer with 32 units (with relu activation). End the network with a softmax layer. {`Sequential, Add, Conv2D, MaxPooling2D, Flatten, Dense from keras`}

2. Get the summary of your network to check it is correct.{`summary from keras`}

3. Fit the CNN to your training and validation data. Use the same loss function, batch size, optimizer and Early Stopping callback that were used for the MLP (warning: training may take more than 10 minutes if you do not have a GPU).
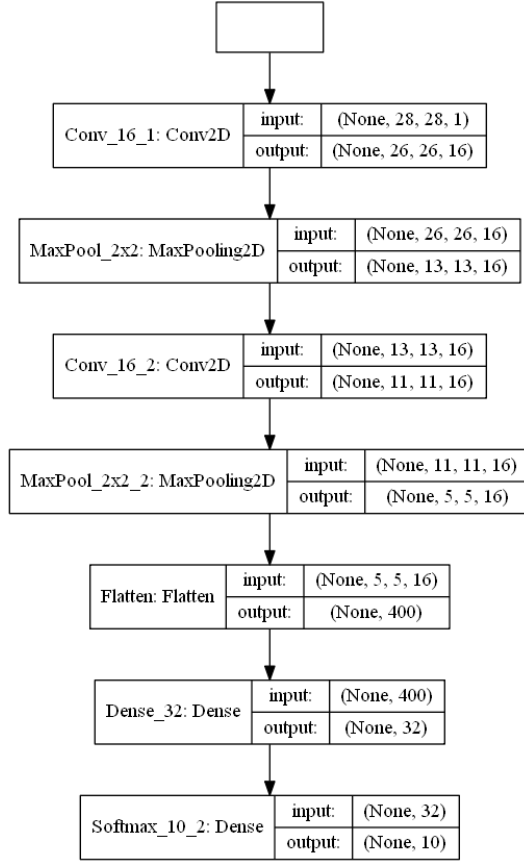
Figure 3: Graph plot of the CNN model.

4. Plot the evolution of the training loss and the validation loss

5. Evaluate performance (accuracy and confusion matrix) on the test data

6. Visualize the feature maps obtained at the output of both convolutional layers, using the function that was provided, for several test images of your choice. {`visualize_activations`}

1.4   Comments

1. Comment on the existence of overfitting in the models that you trained.

2. Comment on the differences between MLP and CNN in what regards performance and number of parameters.

3. Comment on the confusion matrix results for the CNN.

4. Comment on the activation images that are obtained.